

[Table of contents \(../toc.ipynb\)](#)

NumPy



- Numpy is probably the most applied and used Python package. Almost any scientific package builds on numpy.
- Numpy provides multidimensional array objects, and allows to deal with matrix computations.
- Add to this, numpy provides interfaces to C and C++.
- Numpy's methods are very efficient implemented.
- Please find numpy's documentation here <https://numpy.org/> (<https://numpy.org/>).
- Numpy is one building block of Python's scientific eco system. If you want to know more about scientific Python, consult the scipy lecture notes on <https://scipy-lectures.org/> (<https://scipy-lectures.org/>).

Numpy import

First, let us import numpy, create a vector and compare this with a Python list, which ships per default with Python.

```
In [1]: import numpy as np
```

```
In [2]: a_list = range(0, 8000)  
a_np_array = np.arange(0, 8000)
```

Now, let us compare how long it takes to add 3 to each element of the list and the numpy array with `%timeit` magic command.

```
In [3]: %timeit [a_i + 3 for a_i in a_list]
```

412 μ s \pm 16.3 μ s per loop (mean \pm std. dev. of 7 runs, 1000 loops each)

```
In [4]: %timeit a_np_array + 3
```

2.67 μ s \pm 40.4 ns per loop (mean \pm std. dev. of 7 runs, 100000 loops each)

Numpy is much faster than the list and the code of numpy is easier to read!

Create arrays manually

- The "manual" syntax to create a numpy array is
 - `np.array([x_0, x_1, ..., x_n])` for one dimensional arrays
 - and `np.array([[x_0, x_1, ... x_n], [y_0, y_1, ..., y_n]])` for multidimensional arrays.
- Add to this, various functions like `np.ones()`, `np.eye()`, `np.arange()`, `np.linspace()`, and many more create specific arrays which are often required in matrix computing.

```
In [5]: # a one dimensional array  
a = np.array([0, 1, 2, 3])  
a
```

```
Out[5]: array([0, 1, 2, 3])
```

```
In [6]: # here two dimesions  
b = np.array([[0, 1], [2, 3]])  
b
```

```
Out[6]: array([[0, 1],  
               [2, 3]])
```

```
In [7]: # now check their shapes  
print(a.shape)  
print(b.shape)
```

```
(4,)
```

```
(2, 2)
```

Basic attributes

In addition to `ndarray.shape`, numpy arrays contain the attributes:

- `ndarray.ndim` which is the dimension of the array,
- `ndarray.size` which is the number of elements in the array,
- `ndarray.dtype` which is the type of the array (int16, int32, uint16, ..., the default is int64, or float64).

```
In [8]: print('a.dtype = ', a.dtype)
        print('b.dtype = ', b.dtype)
```

```
a.dtype = int64
b.dtype = int64
```

```
In [9]: print('a.ndim = ', a.ndim)
        print('b.ndim = ', b.ndim)
```

```
a.ndim = 1
b.ndim = 2
```

```
In [10]: print('a.size = ', a.size)
         print('b.size = ', b.size)
```

```
a.size = 4
b.size = 4
```

You can also specify the type.

```
In [11]: np.array([1, 2, 3], dtype=np.uint16)
```

```
Out[11]: array([1, 2, 3], dtype=uint16)
```


More array constructors

np.arange

- Create arrays with start, stop, and step size.

```
In [12]: # linear growing array, zero based  
np.arange(8)
```

```
Out[12]: array([0, 1, 2, 3, 4, 5, 6, 7])
```

```
In [13]: # np.arange with start, end, step call  
np.arange(0, 12, 2)
```

```
Out[13]: array([ 0,  2,  4,  6,  8, 10])
```

np.linspace

- Create equidistant arrays with start, stop, and number of points.

```
In [14]: np.linspace(0, 10, 3)
```

```
Out[14]: array([ 0.,  5., 10.])
```

```
In [15]: np.linspace(0, 10, 13)
```

```
Out[15]: array([ 0.          ,  0.83333333,  1.66666667,  2.5          ,  3.33333333,  
                4.16666667,  5.          ,  5.83333333,  6.66666667,  7.5          ,  
                8.33333333,  9.16666667, 10.          ])
```

Special matrices

```
In [16]: np.ones(3)
```

```
Out[16]: array([1., 1., 1.])
```

```
In [17]: np.ones(shape=(3, 3))
```

```
Out[17]: array([[1., 1., 1.],  
                [1., 1., 1.],  
                [1., 1., 1.]])
```

```
In [18]: np.eye(3)
```

```
Out[18]: array([[1., 0., 0.],  
               [0., 1., 0.],  
               [0., 0., 1.]])
```

```
In [19]: np.zeros(3)
```

```
Out[19]: array([0., 0., 0.])
```

```
In [20]: np.zeros(shape=(3, 3))
```

```
Out[20]: array([[0., 0., 0.],  
                [0., 0., 0.],  
                [0., 0., 0.]])
```

Random number arrays

- Numpy's `np.random.xxx` module offers many random number generators.

```
In [21]: np.random.rand(3, 3) # uniform distribution over [0, 1)
```

```
Out[21]: array([[0.85364355, 0.20699322, 0.18675912],  
               [0.98231555, 0.52772497, 0.54494537],  
               [0.93818789, 0.91388519, 0.44828561]])
```

```
In [22]: np.random.randn(3, 3) # standard normal distribution
```

```
Out[22]: array([[-2.20626319, -0.63197416,  0.80843017],  
               [-0.33016822,  0.56584215, -0.16235999],  
               [ 1.76202852, -0.18058352, -0.03921261]])
```

Operators

- All operators $+$, $-$, $*$, $>$, ..., work element wise per default.
- A new array will be created unless you use the $+=$, $-=$ and so forth operators.
- Matrix multiplication can be done with $@$ operator (requires Python 3.5) or `.dot` method.

```
In [23]: a = np.ones(4)
         a + 4
```

```
Out[23]: array([5., 5., 5., 5.])
```

```
In [24]: a * 2
```

```
Out[24]: array([2., 2., 2., 2.])
```

```
In [25]: a[2] = 5
         a > 4
```

```
Out[25]: array([False, False,  True, False])
```

```
In [26]: a = np.array([0, 1, 2, 3])  
        b = np.array([0, 1, 2, 3])  
        a * b
```

```
Out[26]: array([0, 1, 4, 9])
```

```
In [27]: a @ b # matrix product
```

```
Out[27]: 14
```

```
In [28]: a.dot(b) # the "old" way to write a matrix product
```

```
Out[28]: 14
```


Universal functions

- Numpy offers almost all mathematical functions you might need, such as
 - `exp`
 - `max`, `min`
 - `sqrt`
 - `argmax`
 - `median`, `mean`, `stdev`
 - ...

```
In [29]: a = np.linspace(0, 8, 16)  
np.exp(a)
```

```
Out[29]: array([1.00000000e+00, 1.70460487e+00, 2.90567775e+00, 4.95303242e+00,  
8.44296317e+00, 1.43919161e+01, 2.45325302e+01, 4.18182703e+01,  
7.12836271e+01, 1.21510418e+02, 2.07127249e+02, 3.53070116e+02,  
6.01845038e+02, 1.02590798e+03, 1.74876773e+03, 2.98095799e+03])
```

Shape modification

- There are many ways to change the shape of an array.
- Most prominent methods are `reshape` and `transpose`.

```
In [30]: a = np.arange(12).reshape(3, 4)
a
```

```
Out[30]: array([[ 0,  1,  2,  3],
               [ 4,  5,  6,  7],
               [ 8,  9, 10, 11]])
```

```
In [32]: a.reshape(6, 2)
```

```
Out[32]: array([[ 0,  1],
               [ 2,  3],
               [ 4,  5],
               [ 6,  7],
               [ 8,  9],
               [10, 11]])
```

```
In [33]: a.reshape(a.size) # Flatten an array, similar as a.shape(-1), a.flatten() or a.ravel()
```

```
Out[33]: array([ 0,  1,  2,  3,  4,  5,  6,  7,  8,  9, 10, 11])
```

```
In [32]: a = np.zeros((1, 4))  
a
```

```
Out[32]: array([[0., 0., 0., 0.]])
```

```
In [33]: a.transpose() # transpose
```

```
Out[33]: array([[0.],  
                [0.],  
                [0.],  
                [0.]])
```

```
In [34]: a.T # transpose, short hand code
```

```
Out[34]: array([[0.],  
               [0.],  
               [0.],  
               [0.]])
```

Indexing and iterating

- Is very similar to lists.
- Indexing is done by braces `[]`.

```
In [36]: a = np.random.rand(3, 3)
```

```
In [37]: # modify one element  
a[1, 1] = 96
```

```
In [39]: a[0, :]
```

```
Out[39]: array([0.90156267, 0.15860076, 0.92586081])
```

```
In [40]: # iteration over first dimension
```

```
for row in a:  
    print(row)
```

```
[0.90156267 0.15860076 0.92586081]
```

```
[ 0.82601745 96.          0.33312838]
```

```
[0.60133987 0.53074791 0.81656546]
```


If you want to iterate over all elements instead, use the flat attribute.

```
In [41]: for a_i in a.flat:  
         print(a_i)
```

```
0.9015626721609997  
0.15860075987052058  
0.9258608073857245  
0.8260174527163244  
96.0  
0.3331283759775331  
0.6013398695667754  
0.5307479085804412  
0.8165654631841754
```

Slicing arrays

- Slicing is also done with braces.
- There is a great "conversion" table for Matlab users here <https://numpy.org/devdocs/user/numpy-for-matlab-users.html> (<https://numpy.org/devdocs/user/numpy-for-matlab-users.html>).
- The basic slicing syntax is `[start:stop:step]`, see a full tutorial here <https://docs.scipy.org/doc/numpy/reference/arrays.indexing.html> (<https://docs.scipy.org/doc/numpy/reference/arrays.indexing.html>).

```
In [42]: a = np.arange(20)
a
```

```
Out[42]: array([ 0,  1,  2,  3,  4,  5,  6,  7,  8,  9, 10, 11, 12, 13, 14, 15, 16,
                17, 18, 19])
```

All indices are zero based, the stop is not inclusive, and negative means to reverse counting (count from end to start).

```
In [43]: a[-1] # last element
```

```
Out[43]: 19
```

```
In [44]: a[0] # first element
```

```
Out[44]: 0
```

```
In [45]: a[2:-1:2] # start from index 2 to last index and take every second value
```

```
Out[45]: array([ 2,  4,  6,  8, 10, 12, 14, 16, 18])
```

Stacking arrays

```
In [46]: a = np.ones(3)
         b = np.ones(3) + 2
         np.vstack((a, b))
```

```
Out[46]: array([[1., 1., 1.],
               [3., 3., 3.]])
```

```
In [47]: np.hstack((a, b))
```

```
Out[47]: array([1., 1., 1., 3., 3., 3.])
```

Comparing floats

- The `numpy.allclose` is ideal to compare arrays element wise with relative or absolute tolerance.
- The syntax is `np.allclose(a, b, rtol=1e-05, atol=1e-08, equal_nan=False)`.

```
In [48]: np.allclose([1e10, 1e-7], [1.00001e10, 1e-8])
```

```
Out[48]: False
```

```
In [49]: np.allclose([1e10,1e-7], [1.00001e10,1e-8], atol=1e-3) # with larger absolute tolerance, it returns True
```

```
Out[49]: True
```

Linear algebra

- Numpy comes with a linear algebra module `numpy.linalg`.
- Next comes an example to solve an overdetermined system of equations where:
 - $A \in \mathbb{R}^{n \times m}$ denotes input data,
 - $X \in \mathbb{R}^n$ is the parameter vector,
 - and $b \in \mathbb{R}^m$ is the output vector which follows $b = A^\top X$

```
In [50]: n = 2; m = 20
A = np.random.rand(n, m) # input data, two dimensions, 20 samples
X = np.array([[3], [5]]) # two dimensional parameter vector
b = A.T @ X # model the output

# now solve the system of equations
np.linalg.solve(a=(A @ A.T), b=(A @ b))
```

```
Out[50]: array([[3.],
               [5.]])
```

```
In [51]: np.linalg.inv(A @ A.T)
```

```
Out[51]: array([[ 0.34901833, -0.31003372],
               [-0.31003372,  0.46588732]])
```

Read csv data

- Numpy's `genfromtxt` is very convenient to read csv files into numpy arrays.
- The syntax is `dat = genfromtxt('my_file.csv', delimiter=',')`, and there are many additional import properties like `skip_rows`, `missing_values`, ...

We will load this .csv file into a numpy array.

```
"Time", "Torque"  
0, 200  
1, 220  
2, 225  
3, 230  
4, 231
```



```
In [81]: # This if else is a fix to make the file available for Jupyter and Travis CI  
import os  
  
if os.path.isfile('my_file.csv'):  
    file = 'my_file.csv'  
else:  
    file = '02_tools-and-packages/my_file.csv'
```

```
In [82]: from numpy import genfromtxt  
  
genfromtxt(file, delimiter=',', skip_header=1)
```

```
Out[82]: array([[ 0., 200.],  
                [ 1., 220.],  
                [ 2., 225.],  
                [ 3., 230.],  
                [ 4., 231.]])
```

You can also use the column names during the import with `names=True`.

```
In [53]: dat = genfromtxt(file, delimiter=',', names=True)
         dat
```

```
Out[53]: array([(0., 200.), (1., 220.), (2., 225.), (3., 230.), (4., 231.)],
              dtype=[('Time', '<f8'), ('Torque', '<f8')])
```

```
In [54]: print(dat["Time"])
         print(dat["Torque"])
```

```
[0. 1. 2. 3. 4.]
[200. 220. 225. 230. 231.]
```

Matrix computations become so much simpler with numpy

Remember from last lesson that we used lists to plot for instance a parabola. This code was not that handy.

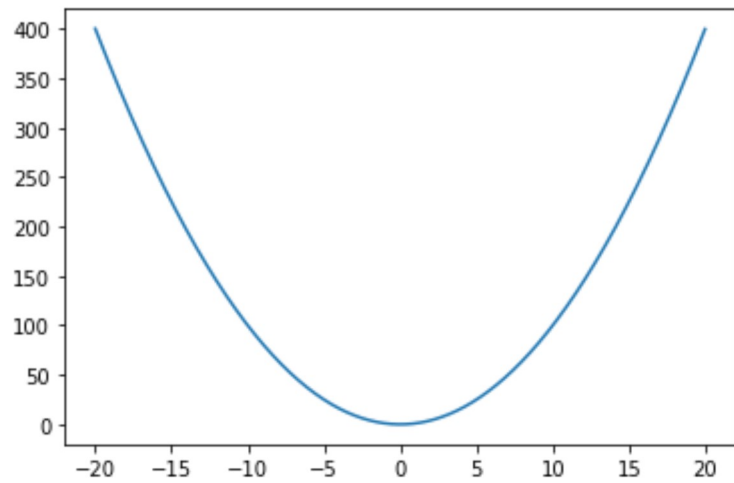
```
x = [i for i in range(-20, 21)]  
y = [x_i**2 for x_i in x]
```

With numpy, such tasks become very simple.

```
In [55]: %matplotlib inline
import matplotlib.pyplot as plt

x = np.arange(-20, 20, 0.01)

plt.plot(x, x**2) # x**2 is easy to read compared with a list [x_i**2 for x_i in x]
plt.show()
```



Exercise: Numpy mini project (20 minutes)



Now that you are familiar with matplotlib and numpy, you can solve the first more elaborate data task.

This is what you should strive for:

- Write a Python script which creates a linear multiplier model of type $AX \approx b$. The dimension of A should be 1000 times 2, and the respective dimension of X becomes 2.
- Select two values for X as you like. These are the true model parameters.
- Generate input data (random numbers or sine waves, as you like) and compute b .
- The \approx sign in the above equation is due to noise that you should add to b .
- The noise should be Gaussian $\mathcal{N} \sim (0, 0.01)$.
- Plot the noisy data b .
- Use `np.linalg.solve` to compute the least squares solution \hat{X} for the parameters.
- Use the same input data and \hat{X} to compute and to plot the fit of the least squares solution.
- Print the true and estimated parameters.

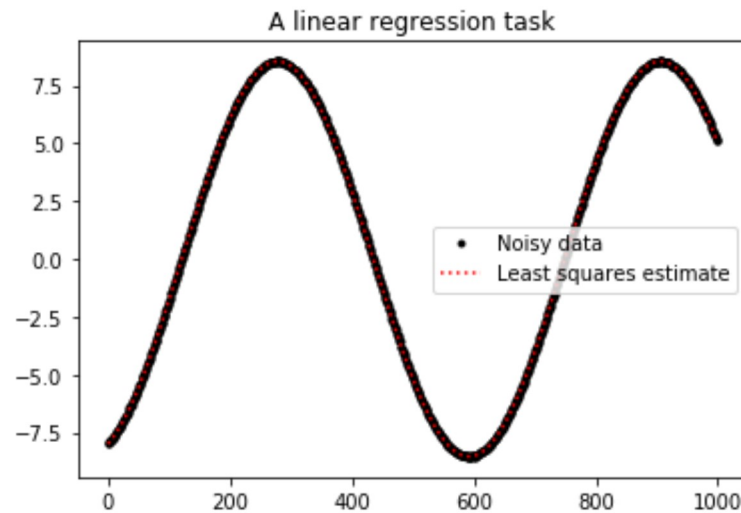
Solution

Please find one possible solution in [solution_numpy.py](#) ([solution_numpy.py](#)) file.

```
In [56]: %run solution_numpy.py
```

```
True params [ 3. -8.]
```

```
Estimated params [ 2.99942396 -7.99943241]
```



<Figure size 432x288 with 0 Axes>